

# Understanding the .bss Segment in C Programming

How Uninitialized Variables Save Executable Space



MOHIT

JAN 19, 2025



20



Sl

## Introduction

One detail that often puzzles programmers, particularly those new to systems programming, is the purpose and necessity of the `.bss` segment. This blog post aims to understand the `.bss` segment, explaining why it is required, how it interacts with other memory segments, and how it affects the performance and efficiency of your programs.

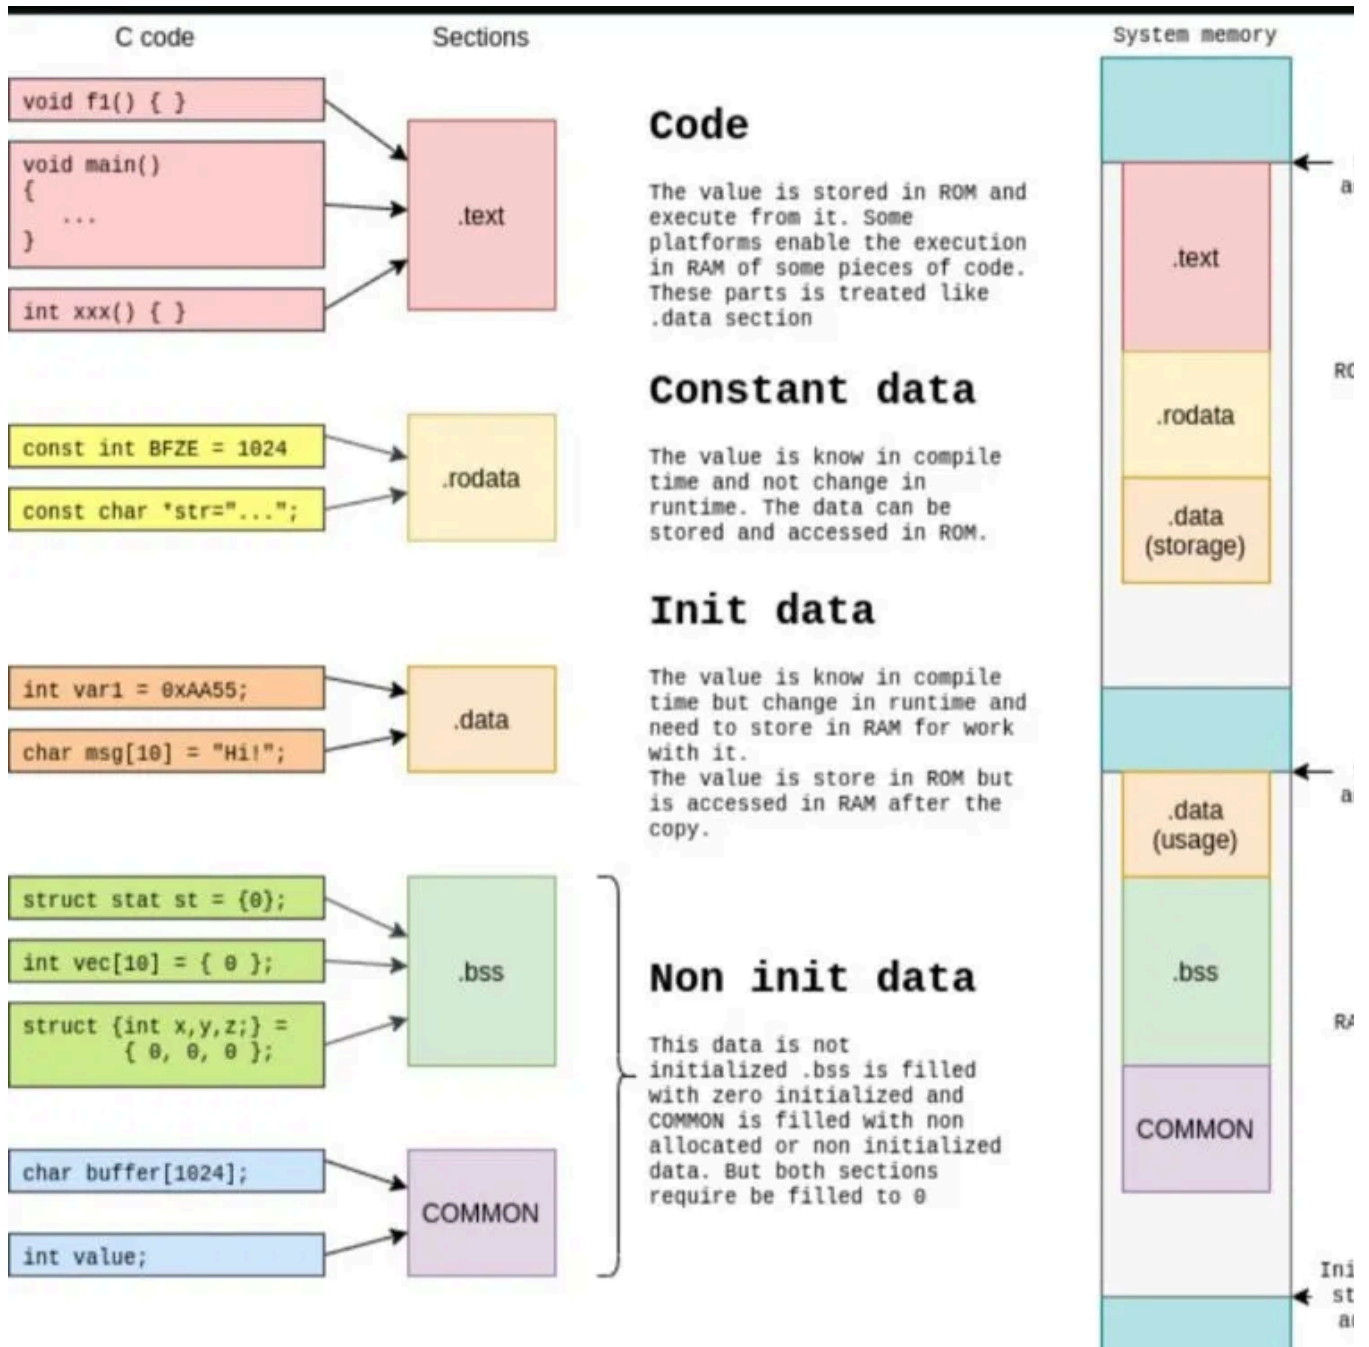
## Memory Segmentation in C

Before diving into the specifics of the `.bss` segment, it's essential to have a basic understanding of memory segmentation in C programming. When a C program is executed, its memory is divided into several segments, each serving a distinct purpose.

1. **Text (Code) Segment:** Contains the executable instructions of the program.
2. **Data Segment:** Holds global and static variables that are initialized by the programmer.
3. **.bss Segment:** Contains global and static variables that are uninitialized or initialized to zero.
4. **Heap:** Dynamic memory allocated during program execution using functions `malloc()`.
5. **Stack:** Used for local variables and function call management.

Thanks for reading Low-Level Lore! Subscribe for free to receive new posts and support my work.

This post focuses on the `.data` and `.bss` segments, as they are closely related and often sources of confusion.



## The Role of the .bss Segment

# What is the .bss Segment?

The `.bss` segment is a part of a program's memory that is used to store global and static variables that are uninitialized or initialized to zero. The name `.bss` is derived from the older assembler keyword "block started by symbol."

## Why Do We Need a Separate .bss Segment?

The primary reason for having a separate `.bss` segment is **memory optimization**. When a program is stored on disk, the `.data` segment contains the values of initialized variables, while the `.bss` segment only records the amount of memory needed for uninitialized variables. This distinction significantly reduces the size of the executable file because the `.bss` segment does not store any actual data—just the information about how much memory to allocate.

## Example Program 1: Demonstrating .data and .bss Segments

Let's consider the following program:

```
#include <stdio.h>
#include <stdlib.h>

int a[10] = { 1, 2, 3, 4, 5, 6, 7, 8, 9 };
int b[20]; /* Uninitialized, so in the .bss segment */

int main()
{
    ;
}
```

### Explanation:

- **Array a:** Initialized with values, so it resides in the `.data` segment.
- **Array b:** Uninitialized, so it resides in the `.bss` segment.

## How to Run and Inspect the Segments:

### 1. Compile the Program:

Use the `-Wall` flag for warnings and `-g` for debugging symbols.

```
gcc -Wall -g example1.c -o example1
```

### 1. Inspect the Executable:

Use the `readelf` or `objdump` tool to view the segments.

```
readelf -a example1 | less
```

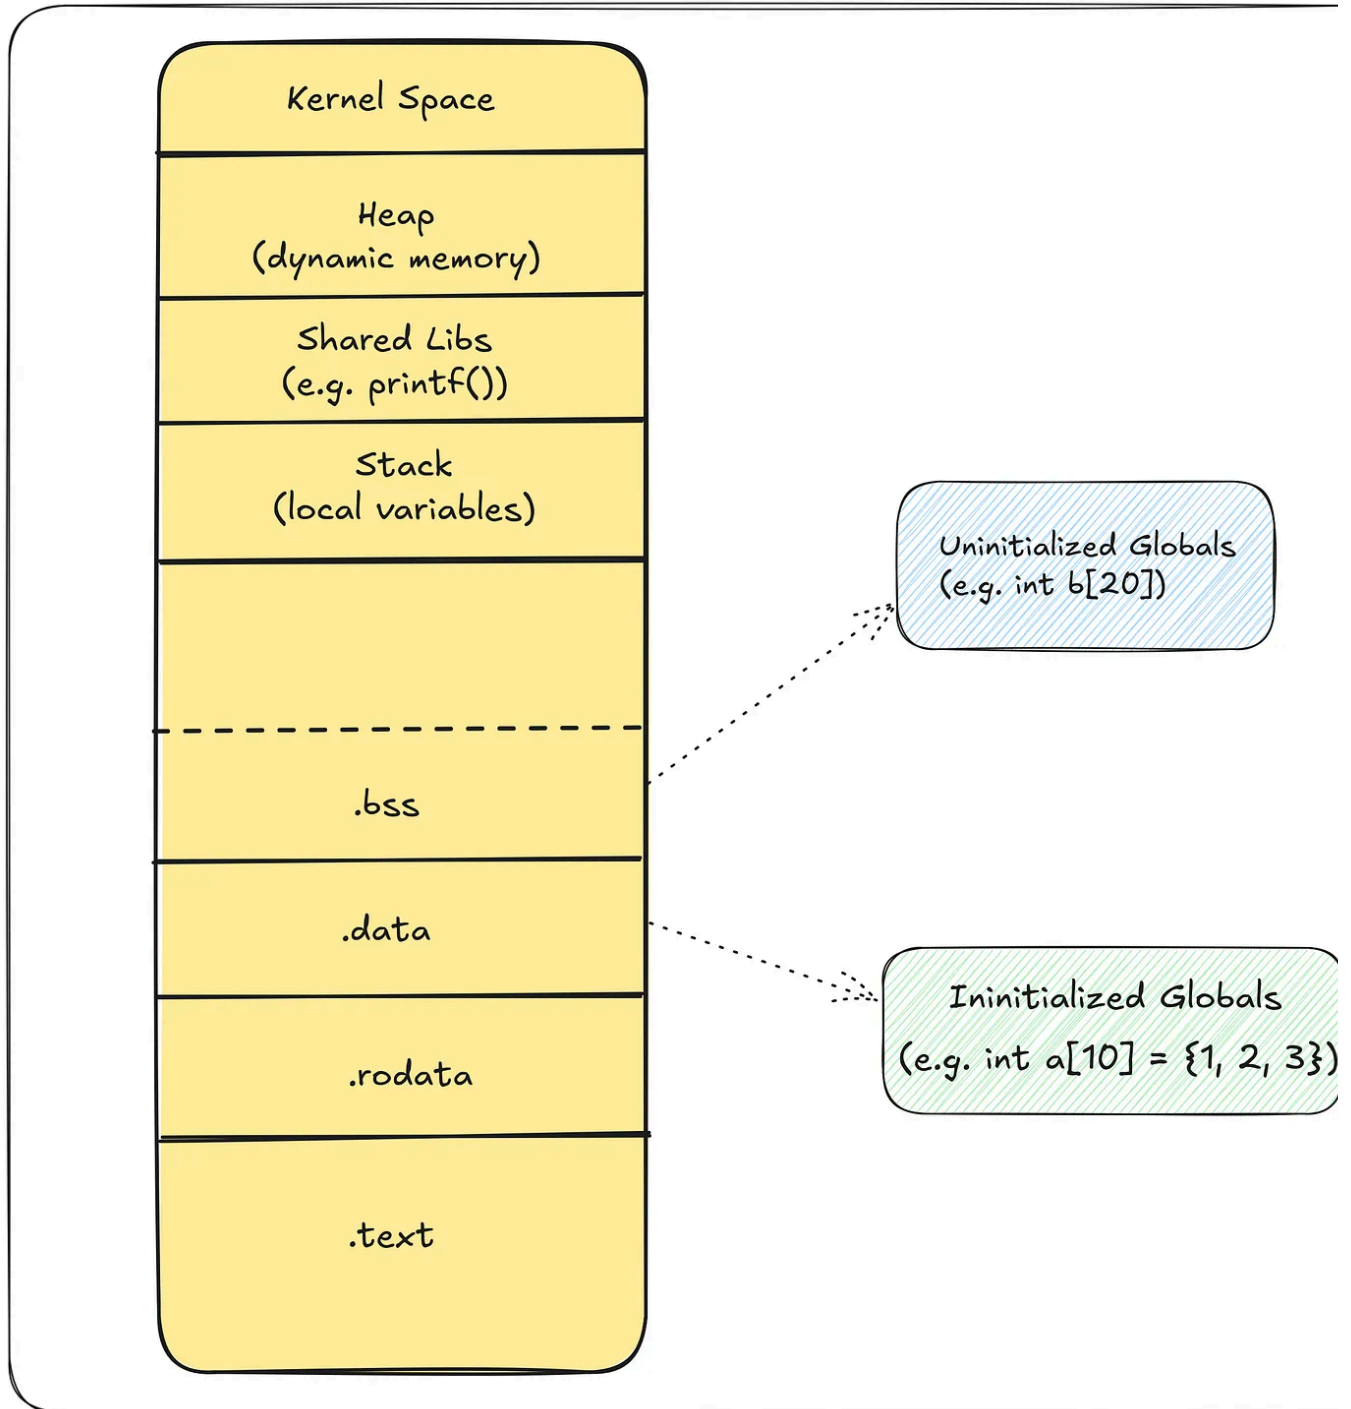
1. Look for the `.data` and `.bss` sections to see their sizes and contents.

## Expected Output:

- The `.data` segment will contain the initialized values of the array `a`.
- The `.bss` segment will reserve space for an array `b` but will not contain any data in the executable file.

## What to Understand:

- The `.bss` segment does not consume space in the executable file, which keeps file size small.
- At runtime, the operating system allocates the necessary memory for the `.bss` segment and initializes it to zero.



## Runtime Initialization of Variables

### Example Program 2: Initializing a .bss Variable at Runtime

Consider the following program:

```
#include <stdio.h>
#include <stdlib.h>
```

```
int var[10]; /* Uninitialized, so in the .bss segment */

int main()
{
    var[0] = 20; /* Initializing at runtime */
    printf("var[0] = %d\n", var[0]);
    return 0;
}
```

### Explanation:

- **Array var:** Declared but not initialized, so it resides in the **.bss** segment.
- **Initialization in main():** The array is initialized at runtime.

### How to Run and Observe Behavior:

#### 1. Compile the Program:

```
gcc -Wall -g example2.c -o example2
```

#### 1. Run the Program:

```
./example2
```

#### 1. Inspect the Output:

```
var[0] = 20
```

### What to Understand:

- Even though **var** is initialized at runtime, it still resides in the **.bss** segment

- The `.bss` segment is zero-initialized at program startup, and any further initialization happens at runtime.
- The location of `var` in memory does not change; it remains in the `.bss` segment throughout the program's execution.

## Program Startup and Memory Initialization

### The Program's Startup Sequence

When a C program starts, the operating system performs several steps to prepare the program's memory:

1. **Allocate Memory for the Program:**

- Memory is allocated to the `.text`, `.data`, `.bss`, heap, and stack segments.

2. **Initialize the `.data` Segment:**

- The initialized values from the executable file are copied into the `.data` segment in memory.

3. **Zero-Initialize the `.bss` Segment:**

- The memory allocated for the `.bss` segment is set to zero.

4. **Call the `main()` Function:**

- The program begins execution from `main()`.

### The Impact of Initialization on Executable Size

If all variables, whether initialized or not, were placed in the `.data` segment, the executable file would be significantly larger. This is because the `.data` segment would contain a lot of unnecessary zeros for uninitialized variables.

By separating uninitialized variables into `.bss` segments, the executable file size is reduced, and memory is used more efficiently.

### Example Program 3: Comparing Executable Sizes

Consider two versions of a program: one where variables are uninitialized and another where they are initialized to zero.

### Version 1: Uninitialized Variables

```
#include <stdio.h>

int a[1000000]; /* Uninitialized, so in the .bss segment */

int main()
{
    printf("Size of array a: %lu bytes\n", sizeof(a));
    return 0;
}
```

### Version 2: Initialized Variables

```
#include <stdio.h>

int a[1000000] = {1}; /* Initialized, so in the .data segment */

int main()
{
    printf("Size of array a: %lu bytes\n", sizeof(a));
    return 0;
}
```

### How to Run and Compare Sizes:

#### 1. Compile Both Programs:

```
gcc -Wall -O2 version1.c -o version1
gcc -Wall -O2 version2.c -o version2
```

#### 1. Check Executable Sizes:



```
ls -lh version1 version2
```

Output:

```
→ ~ ls -lh version1 version2
-rwxrwxr-x 1 chessman chessman 16K Jan 19 11:29 version1
-rwxrwxr-x 1 chessman chessman 3.9M Jan 19 11:29 version2
→ ~ size version1 version2
  text    data    bss      dec     hex filename
  1402      600 4000032 4002034 3d10f2 version1
  1402 4000616      8 4002026 3d10ea version2
```

## Explanation:

- **version1:** The executable size is small (16.0K) because the `.bss` segment does not store data in the file. It only reserves space at runtime.
- **version2:** The executable size is large (3.9M) because the `.data` segment stores the initialized array in the file.

## What to Understand:

- Placing variables in the `.bss` segment instead of `.data` reduces the executable size on disk.
- This optimization is crucial in environments with limited storage, such as embedded systems.

# Practical Implications and Best Practices

## Understanding Variable Placement

- **Initialized Variables:** Place in `.data` segment.
- **Uninitialized Variables:** Place in `.bss` segment.

## Memory Efficiency

- Avoid initializing variables to zero in the source code if they can be left uninitialized and rely on the `.bss` segment's zero-initialization.
- This practice reduces executable size and improves load times.

## Debugging and Memory Analysis

- Use tools like `readelf`, `objdump`, or memory profilers to inspect how variables are placed in memory.
- Understanding memory layout can help in debugging issues related to memory corruption or leaks.

## Embedded Systems Considerations

- In embedded systems, where ROM (Read-Only Memory) is limited, the `.bss` segment's optimization is vital.
- Reducing the `.data` segment's size by using the `.bss` segment can save precious ROM space.

## Conclusion

The `.bss` segment plays a critical role in optimizing memory usage and reducing the size of executable files. By separating uninitialized variables from initialized ones, the `.bss` segment ensures that programs are efficient in both memory and disk space. Understanding how and why the `.bss` segment is used can lead to better programming practices, especially in resource-constrained environments.

## Further Reading and Resources

- **ELF (Executable and Linkable Format):** Learn more about how executables are structured.
- **Memory Management in C:** Explore how memory is allocated and managed in programs.

- **Embedded Systems Programming:** Delve into the specifics of programming for systems with limited resources.

## FAQs

### Q1: Can variables in the .bss segment be initialized to values other than zero at runtime?

A1: Yes, variables in the .bss segment are zero-initialized at program startup, but they can be assigned any value at runtime just like variables in the .data segment.

### Q2: Does the .bss segment exist in all operating systems?

A2: Yes, the .bss segment is a standard part of the memory layout in most operating systems that use the ELF format for executables, including Linux and macOS.

### Q3: How can I check which variables are in the .bss segment?

A3: You can use tools like `readelf` or `objdump` to inspect the memory segments of an executable and see which variables are placed in the .bss segment.

### Q4: What happens if I initialize a variable in the .bss segment to a non-zero value?

A4: If you initialize a variable to a non-zero value, the compiler will place it in the .data segment instead of the .bss segment.

## Final Thoughts

Understanding the intricacies of memory management, including the role of the .bss segment, is essential for writing efficient and optimized code. By leveraging the .bss segment effectively, developers can create programs that are not only faster but also more adaptable to various environments, from desktop applications to embedded systems.

Thanks for reading Low-Level Lore! Subscribe for free to receive new posts and support my work.

Thanks for reading Low-Level Lore! This post is public so feel free to share it.



20 Likes

## Discussion about this post

Comments

Restacks



Write a comment...

---

© 2025 Mohit · [Privacy](#) · [Terms](#) · [Collection notice](#)  
[Substack](#) is the home for great culture